
Estimating a Structural Model of Herd Behavior in Financial Markets

Author(s): Marco Cipriani and Antonio Guarino

Source: *The American Economic Review*, JANUARY 2014, Vol. 104, No. 1 (JANUARY 2014), pp. 224-251

Published by: American Economic Association

Stable URL: <https://www.jstor.org/stable/42920693>

JSTOR is a not-for-profit service that helps scholars, researchers, and students discover, use, and build upon a wide range of content in a trusted digital archive. We use information technology and tools to increase productivity and facilitate new forms of scholarship. For more information about JSTOR, please contact support@jstor.org.

Your use of the JSTOR archive indicates your acceptance of the Terms & Conditions of Use, available at <https://about.jstor.org/terms>



JSTOR

is collaborating with JSTOR to digitize, preserve and extend access to *The American Economic Review*

Estimating a Structural Model of Herd Behavior in Financial Markets[†]

By MARCO CIPRIANI AND ANTONIO GUARINO*

We develop a new methodology to estimate herd behavior in financial markets. We build a model of informational herding that can be estimated with financial transaction data. In the model, rational herding arises because of information-event uncertainty. We estimate the model using data on a NYSE stock (Ashland Inc.) during 1995. Herding occurs often and is particularly pervasive on some days. On average, the proportion of herd buyers is 2 percent; that of herd sellers is 4 percent. Herding also causes important informational inefficiencies in the market, amounting, on average, to 4 percent of the asset's expected value. (JEL C58, D82, D83, G12, G14)

In recent years there has been much interest in herd behavior in financial markets. This interest has led researchers to look for theoretical explanations and empirical evidence of herding. There has been, however, a substantial disconnect between the empirical and theoretical literatures: the theoretical work has identified motives for herding in abstract models that cannot easily be brought to the data; the empirical literature has mainly looked for atheoretical, statistical evidence of trade clustering, which is interpreted as herding.

This article takes a novel approach: we develop a theoretical model of herding in financial markets that can be estimated with financial market transaction data. This methodology allows us to measure the quantitative importance of herding, to identify when it happens, and to assess the informational inefficiency that it generates.

*Cipriani: New York Federal Reserve Bank, 33 Liberty Street, New York, NY 10045 (e-mail: marco.cipriani@ny.frb.org); Guarino: Department of Economics, University College London, Gower Street, London WC1E 6BT, UK (e-mail: a.guarino@ucl.ac.uk). We thank Orazio Attanasio, Richard Blundell, Pedro Carneiro, Andrew Chesher, Burkhard Drees, David Easley, Robert Engle, Christopher Flinn, Ana Fostel, Douglas Gale, Toru Kitagawa, Mico Loretan, Lars Nesheim, Maureen O'Hara, Sam Ouliaris, Henri Pages, Nicola Pavoni, Jean Marc Robin, Ling Hui Tan, and seminar participants in various institutions for comments and suggestions. We are grateful to three anonymous referees for providing us with valuable comments and suggestions. We thank Adam Biesenbach, Silvia Camussi, Minyu Chen, Mateusz Giezek, Thibaut Lamadon, Sean Myers, Yuki Sato, Federico Tagliati, and Andreas Uthemann for excellent research assistance. We gratefully acknowledge the financial support of the Fondation Banque de France, the ERC, and the ESRC. This work was partially conducted while Guarino was a visiting scholar at the IMF Institute; Guarino is grateful to the IMF Institute for its hospitality. The views expressed herein are those of the authors and should not be attributed to the IMF, its executive board, or its management; the Federal Reserve Bank of New York, or the Federal Reserve System. We are responsible for any errors.

[†]Go to <http://dx.doi.org/10.1257/aer.104.1.224> to visit the article page for additional materials and author disclosure statement(s).

The theoretical work on herd behavior started with the seminal papers of Banerjee (1992); Bikhchandani, Hirshleifer, and Welch (1992); and Welch (1992).¹ These papers model herd behavior in an abstract environment in which agents with private information make their decisions in sequence. They show that, after a finite number of agents have chosen the same action, all following agents disregard their own private information and imitate their predecessors. More recently, a number of papers (see, among others, Avery and Zemsky 1998; Lee 1998; and Cipriani and Guarino 2008) have focused on herd behavior in financial markets. In particular, these studies analyze a market where informed and uninformed traders sequentially trade a security of unknown value. The price of the security is set by a market maker according to the order flow. The presence of a price mechanism makes it more difficult for herding to arise. Nevertheless, there are cases in which it occurs. In Avery and Zemsky (1998), for instance, herd behavior can occur when there is uncertainty not only about the value of the asset but also about the occurrence of an information event or about the model parameters.

As mentioned above, whereas the theoretical research has tried to identify the mechanisms through which herd behavior can arise, the empirical literature has followed a different track. The existing work (see, e.g., Lakonishok, Shleifer, and Vishny 1992; Grinblatt, Titman, and Wermers 1995; and Wermers 1999) does not test the theoretical herding models directly but analyzes the presence of herding in financial markets through statistical measures of clustering.² These papers find that, in some markets, fund managers tend to cluster their investment decisions more than would be expected if they acted independently. This empirical research on herding is important, as it sheds light on the behavior of financial market participants and in particular on whether they act in a coordinated fashion. As the authors themselves emphasize, however, decision clustering may or may not be due to herding (for instance, it may be the result of a common reaction to public announcements). These papers cannot distinguish spurious herding from true herd behavior, that is, the decision to disregard one's private information to follow the behavior of others (see Bikhchandani and Sharma 2000; and Hirshleifer and Teoh 2009).

Testing models of informational herd behavior is difficult. In such models, a trader herds if he trades against his own private information. The problem that empiricists face is that there are no data on the private information available to traders and, therefore, it is difficult to know when traders decide not to follow it. Our purpose in this paper is to present a methodology to overcome this problem. We develop a theoretical model of herding and estimate it using financial market transaction data. We are able to identify the periods in the trading day in which traders act as herders and to measure the informational inefficiency that this generates. This is the first paper on informational herding that, instead of using a statistical, atheoretical approach, brings a theoretical social learning model to the field data.³

¹ We study only informational herding. Therefore, we do not discuss herd behavior due to reputational concerns or payoff externalities. For an early critical assessment of the literature on herd behavior see Gale (1996). For recent surveys of herding in financial markets see Bikhchandani and Sharma (2000), Vives (2008), and Hirshleifer and Teoh (2009).

² See also the recent paper by Dasgupta, Prat, and Verardo (2011), who study the effect of institutional herding on long-term returns, and the literature cited therein.

³ While there are no direct empirical tests of herding models, there is experimental work that tests these models in the laboratory (see, e.g., Cipriani and Guarino 2005 and 2009; and Drehmann, Oechssler, and Roeder 2005).

Our theoretical analysis builds on the work of Avery and Zemsky (1998), who use a sequential trading model à la Glosten and Milgrom (1985) to study herding in financial markets. Avery and Zemsky (1998) show that, in financial markets, the fact that the price continuously adjusts to the order flow makes herding less likely to arise. However, they also show that herding does occur if there is “event uncertainty” in the market, that is, uncertainty on whether an information event (i.e., a shock to the asset value, on which informed traders receive a signal) has occurred. Since event uncertainty is a typical assumption of sequential trading market microstructure models (starting from Easley and O’Hara 1992), it is a natural way of generating herd behavior in a financial economy.⁴

In our model, herding occurs from a mechanism similar to that exposed by Avery and Zemsky (1998). However, whereas they were interested in providing theoretical examples of herding, our aim is to provide an empirical methodology to gauge the importance of herding in actual financial markets. For this purpose, we build a model of herding that can be estimated with financial market transaction data. In the model, an asset is traded over many days; at the beginning of each day, an informational event may occur that causes the fundamental asset value to change with respect to the previous day. If an informational event has occurred, some traders receive private information on the new asset value.⁵ These traders trade the asset to exploit their informational advantage over the market maker. If no event has occurred, all traders in the market are noise traders, that is, they trade for noninformation reasons only (e.g., liquidity or hedging motives). Whereas the informed traders know that they are in a market with private information (since they themselves are informed), the market maker does not. This asymmetry of information between traders and the market maker implies that the market maker updates the price more “slowly” than an informed trader would in order to take into account the possibility that the asset value may not have changed (in which case all trading activity is due to noninformational motives). As a result, after, for instance, a history of buys, a trader, even with a bad signal, may value the asset more than the market maker does. He will, therefore, trade against his own private information and herd-buy.

We estimate the model with stock market transaction data via maximum likelihood, using a strategy first proposed by Easley, Kiefer, and O’Hara (1997a) to estimate the parameters of the Glosten and Milgrom (1985) model. There is an important difference, however, between Easley, Kiefer, and O’Hara’s (1997a) methodology and ours. In their setup, informed traders are perfectly informed about the value of the asset; as a result, their decisions are never affected by the decisions of previous traders, and they never herd. Therefore, only the total *number* of buys, sells, and no-trades in each day matters; the *sequence* in which these trades arrive is irrelevant. In contrast, in our framework, the precision of private information is one of the parameters that we estimate. This opens the possibility that informed traders may receive noisy signals, and that they may find it optimal to ignore them and engage in herd behavior. In this circumstance, the sequence by which trades arrive in the market does matter: in contrast to Easley, Kiefer, and O’Hara (1997a), we cannot

⁴A similar mechanism is also present in Gervais (1997). Recently, Park and Sabourian (2011) have illustrated the necessary and sufficient conditions on private information for the occurrence of herding.

⁵The event is called informational precisely because some traders in the market receive private information on it.

estimate our model using only the number of buy or sell orders in a given day, but we must consider the whole history of trading activity in each day of trading.

As an illustration of the methodology, we estimate the model using transaction data for a NYSE stock (Ashland Inc.) during 1995. The restriction that private signals are always correct (as in Easley, Kiefer, and O'Hara 1997a) is rejected by the data, which implies both that herd behavior arises in equilibrium and that there is information content in the sequence of trades.

Note that in each day of trading there is always high heterogeneity in trading decisions (i.e., even in days when the fundamental value has increased, we observe many sell orders, and vice versa). If private information were perfectly precise, the only way to account for it would be to have a high proportion of noise traders in the market (indeed, Easley, Kiefer, and O'Hara 1997a estimate that the proportion of noise traders is 83 percent). The advantage of our methodology is that it accounts for the heterogeneity in trading decisions not only through the presence of noise traders, but also by allowing informed traders to receive the wrong piece of information. As one may expect, our estimate of the proportion of informed traders increases substantially with respect to Easley, Kiefer, and O'Hara (1997a); according to our estimates, however, informed traders have a relatively imprecise signal, incorrect 40 percent of the time. In a nutshell, we partially explain the apparent noise in the data as the result of the rational behavior of imperfectly informed traders, as opposed to assuming that it all comes from randomly acting noise traders.

Allowing for an imperfectly precise signal has important consequences for estimates of trading informativeness. A large literature has studied the information content of trading activity using a measure (usually called the PIN, an acronym for private information-based trading) based on the Easley, Kiefer, and O'Hara (1997a) methodology (i.e., assuming that all informed traders receive the correct information). Using that methodology, the measure of information-based activity in our sample would be 9 percent. Using our methodology instead, we obtain 19 percent. The difference is due to the fact that in the previous literature incorrect trades (e.g., selling in a good-event day) can be due only to exogenous, noninformative (e.g., liquidity) reasons, whereas in our setup we do not exclude the possibility that they may be due to informed traders who either receive incorrect information or herd.⁶

Given our estimated parameters, we study how traders' beliefs evolve during each day of trading. By comparing these beliefs to the prices, we are able to identify periods of the trading day in which traders herd. In most of the trading periods, a positive (albeit small) measure of informed traders herd. In an information-event day, on average, between 2 percent (4 percent) of informed traders herd-buy (sell). We also compare our structural estimates of herding with the atheoretical herding measure proposed by Lakonishok, Shleifer, and Vishny (1992). We show that when applied to transaction data, this measure attributes to herding a higher proportion of trades than our structural estimates. The reason is that it does not disentangle the clustering

⁶ As we explain later, in the context of our analysis, the PIN that one would estimate in the standard way (i.e., as in Easley et al. 1996) does not measure the proportion of informed trading activity in the market, but rather the proportion of trading activity stemming from informed traders with the correct signal.

of trading decisions due to true informational herding from clustering due to other reasons (such as the exposure to a common information event).

Herd behavior generates serial dependence in trading patterns, a phenomenon documented in the empirical literature. Herding also causes informational inefficiencies in the market. On average, the misalignment between the price we observe and the price we would observe in the absence of herding is equal to 4 percent of the asset's unconditional fundamental value.

The rest of the article is organized as follows. Section I describes the theoretical model. Section II presents the likelihood function. Section III describes the data. Section IV presents the results. Section V concludes. An online Appendix contains the proofs and other supplementary material.

I. The Model

Following Easley and O'Hara (1987), we generalize the original Glosten and Milgrom (1985) model to an economy where trading happens over many days.

An asset is traded by a sequence of traders who interact with a market maker. Trading days are indexed by $d = 1, 2, 3, \dots$. Time within each day is discrete and indexed by $t = 1, 2, 3, \dots$

The Asset.—We denote the fundamental value of the asset in day d by V_d . The asset value does not change during the day but can change from one day to the next. At the beginning of the day, with probability $1 - \alpha$ the asset value remains the same as in the previous day ($V_d = v_{d-1}$), and with probability α it changes. In each day d , the value of the asset in the previous day $d - 1$, v_{d-1} , is known to all market participants.⁷ As we will explain, when the value of the asset changes from one day to the other, there are informed traders in the market; for this reason, we say that an information event has occurred. If an information event occurs, with probability $1 - \delta$ the asset value decreases to $v_{d-1} - \lambda^L$ ("bad informational event"), and with probability δ it increases to $v_{d-1} + \lambda^H$ ("good informational event"), where $\lambda^L > 0$ and $\lambda^H > 0$. Informational events are independently distributed over the days of trading. To simplify the notation, we define $v_d^H := v_{d-1} + \lambda^H$ and $v_d^L := v_{d-1} - \lambda^L$. Finally, we assume that $(1 - \delta)\lambda^L = \delta\lambda^H$, which, as will become clear later, guarantees that the closing price is a martingale.

The Market.—The asset is exchanged in a specialist market. Its price is set by a market maker who interacts with a sequence of traders. At any time $t = 1, 2, 3, \dots$ during the day a trader is randomly chosen to act and can buy, sell, or decide not to trade. Each trade consists of the exchange of one unit of the asset for cash. The trader's action space is, therefore, $\mathcal{A} = \{\text{buy, sell, no trade}\}$. We denote the action of the trader at time t in day d by X_t^d and the history of trades and prices until time $t - 1$ of day d by H_t^d .

⁷For more comments on this point, see footnote 18. Note that v_{d-1} is the realization of the random variable V_{d-1} . Throughout the text, we will denote random variables with capital letters and their realizations with lowercase letters.

The Market Maker.—At any time t of day d , the market maker sets the prices at which a trader can buy or sell the asset. When posting these prices, the market maker must take into account the possibility of trading with traders who (as we shall see) have some private information on the asset value. He will set different prices for buying and for selling; that is, there will be a bid-ask spread (Glosten and Milgrom 1985). We denote the ask price (the price at which a trader can buy) at time t by a_t^d and the bid price (the price at which a trader can sell) by b_t^d .

Due to (unmodeled) potential competition, the market maker makes zero expected profits by setting the ask and bid prices equal to the expected value of the asset conditional on the information available at time t and on the chosen action, that is,

$$a_t^d = E(V_d | h_t^d, X_t^d = \text{buy}, a_t^d, b_t^d),$$

$$b_t^d = E(V_d | h_t^d, X_t^d = \text{sell}, a_t^d, b_t^d).$$

Following Avery and Zemsky (1998), we will sometimes refer to the market maker's expectation conditional on the history of trades only as the "price" of the asset, and we will denote it by $p_t^d = E(V^d | h_t)$.⁸

The Traders.—There are a countable number of traders. Traders act in an exogenous sequential order. Each trader is chosen to take an action only once, at time t of day d . Traders are of two types, informed and noise. The trader's own type is private information.

In no-event days, all traders in the market are noise traders. In information-event days, at any time t an informed trader is chosen to trade with probability μ and a noise trader with probability $1 - \mu$, with $\mu \in (0, 1)$.⁹

Noise traders trade for unmodeled (e.g., liquidity) reasons: they buy with probability $\frac{\varepsilon}{2}$, sell with probability $\frac{\varepsilon}{2}$, and do not trade with probability $1 - \varepsilon$ (with $0 < \varepsilon < 1$). Informed traders have private information on the asset value. They receive a private signal on the new asset value and observe the previous history of trades and prices, and the current prices.¹⁰ The private signal S_t^d has the following value-contingent densities:

$$g^H(s_t^d | v_d^H) = 1 + \tau(2s_t^d - 1),$$

$$g^L(s_t^d | v_d^L) = 1 - \tau(2s_t^d - 1),$$

with $\tau \in (0, \infty)$. (See Figure 1.)

For $\tau \in (0, 1]$, the support of the densities is $[0, 1]$. In contrast, for $\tau > 1$, the support shrinks to $[\frac{\tau-1}{2\tau}, \frac{\tau-1+2\sqrt{\tau}}{2\tau}]$ for g^H and to $[\frac{\tau+1-2\sqrt{\tau}}{2\tau}, \frac{\tau+1}{2\tau}]$ for g^L (in order for the density functions to integrate to one). Note that, given the value

⁸Standard arguments show that $b_t^d \leq p_t^d \leq a_t^d$ (see Glosten and Milgrom 1985).

⁹In other words, μ should be interpreted as the proportion of informed-based trading decisions in a day (and not as the proportion of informed traders in the population). Of course, in a no-event day, the proportion of informed-based trading decisions is zero.

¹⁰As we will explain later, in the model there is a one-to-one mapping from trades to prices. For this reason, in bringing the model to the data, we need to assume only that traders observe the history of past prices.




FIGURE 1. THE SIGNAL

Note: The figure shows the signal's state-contingent density functions for different values of τ .

of the asset, the signals S_t^d are i.i.d.¹¹ The signals satisfy the monotone likelihood ratio property. At each time t , the likelihood ratio after receiving the signal, $\frac{\Pr(V_d = v_d^H | h_t^d, s_t^d)}{\Pr(V_d = v_d^L | h_t^d, s_t^d)} = \frac{g^H(s_t^d | v_d^H)}{g^L(s_t^d | v_d^L)} \frac{\Pr(V_d = v_d^H | h_t^d)}{\Pr(V_d = v_d^L | h_t^d)}$, is higher than that before receiving the signal if $s_t^d > 0.5$ and lower if $s_t^d < 0.5$. For this reason we refer to a signal larger than 0.5 as a “good signal” and to a signal smaller than 0.5 as a “bad signal.”

The parameter τ measures the informativeness of the signals. When $\tau \rightarrow 0$, the densities are uniform, and the signals are completely uninformative. As τ increases, the signals become more and more informative. For any given $\tau \in (0, 1)$, the support of the distribution of the likelihood ratio is bounded away from 0 and infinity, while for $\tau \geq 1$ it is not. Following Smith and Sørensen (2000), in the first case we say that beliefs are bounded and in the second case that they are unbounded. With bounded beliefs, no signal realizations (even the most extreme ones) reveal the asset value with probability one. With unbounded beliefs, in contrast, some high (low) signal realizations are possible only when the asset value is high (low), and therefore the signal can be perfectly informative.¹² As τ tends to infinity, the measure of perfectly informative signals tends to one.

An informed trader knows that an information event has occurred and that, as a result, the asset value has changed with respect to the previous day. Moreover, his signal is informative on whether the event is good or bad. Nevertheless, according to the signal realization that he receives and the precision τ , he may not be completely sure of the effect of the event on the asset value. For instance, he may know that there has been a change in the investment strategy of a company, but not be sure whether this change will affect the asset value in a positive or negative way. The parameter τ can be interpreted as measuring the precision of the information that the trader receives, or the ability of the trader to process such private information. Finally, note that, given our signal structure, informed traders are heterogeneous,

¹¹ Conditional i.i.d. signals deliver a likelihood function that can be easily brought to the data. Nevertheless, the result that herding arises in equilibrium would also hold in an economy in which signals are not conditionally independent.

¹² In particular, any signal greater than or equal to $\frac{\tau+1}{2\tau}$ reveals that the asset value is v_d^H , whereas a signal lower than or equal to $\frac{\tau-1}{2\tau}$ reveals that the asset value is v_d^L .

since they receive signal realizations with different degrees of informativeness about the asset's fundamental value.

In addition to capturing heterogeneity of information in the market, a linear density function for the signal makes it possible to compute the traders' strategies and the market maker's posted prices analytically. As a result, we obtain a simple and tractable likelihood function. Moreover, in contrast to other specifications such as a discrete signal (e.g., a noisy binary signal), our choice avoids creating a discontinuity in the likelihood function, which would make estimation problematic.

An informed trader's payoff function, $U : \{v_d^L, v_d^H\} \times \mathcal{A} \times [v_d^L, v_d^H]^2 \rightarrow \mathbf{R}^+$, is defined as

$$U(v_d, X_t^d, a_t^d, b_t^d) = \begin{cases} v_d - a_t^d & \text{if } X_t^d = \text{buy}, \\ 0 & \text{if } X_t^d = \text{no trade}, \\ b_t^d - v_d & \text{if } X_t^d = \text{sell}. \end{cases}$$

An informed trader chooses X_t^d to maximize $E(U(V_d, X_t^d, a_t^d, b_t^d) | h_t^d, s_t^d)$ (i.e., he is risk neutral). Therefore, he finds it optimal to buy whenever $E(V_d | h_t^d, s_t^d) > a_t^d$ and to sell whenever $E(V_d | h_t^d, s_t^d) < b_t^d$. He chooses not to trade when $b_t^d < E(V_d | h_t^d, s_t^d) < a_t^d$. Otherwise, he is indifferent between buying and not trading, or selling and not trading.

Note that at each time t , the trading decision of an informed trader can be simply characterized by two thresholds, σ_t^d and β_t^d , satisfying the equalities

$$E[V_d | h_t^d, \sigma_t^d] = b_t^d$$

and

$$E[V_d | h_t^d, \beta_t^d] = a_t^d.$$

An informed trader will sell for any signal realization smaller than σ_t^d and buy for any signal realization greater than β_t^d . Obviously, the thresholds at each time t depend on the history of trades until that time and on the parameter values.¹³

Figure 2 (drawn for the case of a good informational event) illustrates the decision of informed traders. An informed trader buys the asset for a signal higher than the threshold value β_t^d , sells it for a signal lower than σ_t^d , and does not trade otherwise. The measure of informed traders buying or selling is equal to the areas (labeled as "Informed Buy" and "Informed Sell") below the line representing the signal density function.

A. Herd Behavior

To discuss herd behavior, let us start by introducing some formal definitions.

¹³Since noise traders buy and sell with probabilities bounded away from zero, standard arguments prove that the bid and ask prices exist and are unique (see, e.g., Cipriani and Guarino 2008). Similar arguments can be used to prove existence and uniqueness of the thresholds.




FIGURE 2. INFORMED TRADER'S DECISION

Note: The figure illustrates the signal realizations for which an informed trader decides to buy or sell when $V_d = v_d^H$ (the signal density function is conditional on v_d^H).

DEFINITION 1: An informed trader engages in herd buying at time t of day d if
(i) he buys upon receiving a bad signal, that is,

$$E(V_d | h_t^d, s_t^d) > a_t^d \text{ for } s_t^d < 0.5, \text{ and}$$

(ii) the price of the asset is higher than at time 1, that is,

$$p_t^d = E(V_d | h_t^d) > p_1^d = v_{d-1}.$$

Similarly, an informed trader engages in herd selling at time t of day d if

(i) he sells upon receiving a good signal, that is,

$$E(V_d | h_t^d, s_t^d) < b_t^d \text{ for } s_t^d > 0.5, \text{ and}$$

(ii) the price of the asset is lower than at time 1, that is,

$$p_t^d = E(V_d | h_t^d) < p_1^d = v_{d-1}.$$

In other words, a trader herds when he trades against his own private information in order to conform to the information contained in the history of trades, that is, to buy after the price has risen or to sell after the price has fallen. Since traders in our model receive different signals, it may well be (and typically will be the case) that, at a given point in time, traders with less informative signals (i.e., close to 0.5) will herd, whereas traders with more informative signals (close to the extremes of the support) will not. We are interested in periods of the trading day in which traders engage in herd behavior for at least some signal realizations. At any given time t , we can detect whether an informed trader herds for a positive measure of signals by comparing the two thresholds σ_t^d and β_t^d to 0.5. Since a trader engages in herd-buying behavior if he buys despite a bad signal ($s_t^d < 0.5$), there is a positive measure of herd buyers whenever $\beta_t^d < 0.5$.¹⁴ Similarly for herd sellers.

¹⁴We identify an informed trader with the signal he receives: thus, "a positive measure of herd buyers" means "a positive measure of signal realizations for which an informed trader herd buys."

Note that in our model (similarly to Avery and Zemsky's), traders trade against their own private information (i.e., buy after a bad signal or sell after a good one) only to conform to the past trading pattern and never to go against it. Using the language of the social learning literature, in our model agents go against their private information only to herd and never to act as contrarians. This means that whenever condition (i) in Definition 1 is satisfied, so is condition (ii); in other words, condition (ii) is redundant. For instance, an informed trader with, e.g., a bad signal never buys after a history of trades has pushed the price downward with respect to the beginning of the day. Indeed, his expected value after the price has decreased is lower than that of the market maker (because he attaches a higher probability to the event that the sell orders come from informed traders).¹⁵ Therefore, we formally define herd behavior as follows:

DEFINITION 2: *There is herd behavior at time t of day d when there is a positive measure of signal realizations for which an informed trader either herd buys or herd sells, that is, when*

$$\beta_t^d < 0.5 \quad \text{or} \quad \sigma_t^d > 0.5.$$

Figures 3 and 4 show an example of herd buy and herd sell, respectively, in a day with a good information event. The areas below the signal density function and between the thresholds and 0.5 represent the measures of informed traders who herd buy and herd sell.

The reason why herd behavior occurs is that prices move “too slowly” as buy and sell orders arrive in the market. Suppose that, at the beginning of an information-event day, there is a sequence of buy orders. Informed traders, knowing that there has been an information event, attach a certain probability to the fact that these orders come from informed traders with good signals. The market maker, however, attaches a lower probability to this event, as he takes into account the possibility that there was no information event and that all the buys came from noise traders. Therefore, after a sequence of buys, he will update the prices upwards, but by less than the movement in traders' expectations. Because traders and the market maker interpret the history of trades differently, the expectation of a trader with a bad signal may be higher than the ask price, in which case he herd buys (obviously, traders who receive signals close to 0.5 will be more likely to herd, since the history of trades has more weight in forming their beliefs).

We state this result in the next proposition whose proof is in the online Appendix:

PROPOSITION 1: *For any finite τ , herd behavior occurs with positive probability. Furthermore, herd behavior can be misdirected, that is, an informed trader can engage in herd buying (herd selling) in a day when a bad (good) information event has occurred.*

¹⁵The formal proof of the result is contained in the online Appendix.




FIGURE 3. HERD BUY

Note: In the figure, an informed trader buys even after receiving a bad signal (higher than 0.3).




FIGURE 4. HERD SELL

Note: In the figure, an informed trader sells even after receiving a good signal (lower than 0.7).

Avery and Zemsky (1998) have shown how herding can occur because of uncertainty on whether an information event has occurred (see their IS2 information setup). In our model herding arises for the same reason. Our contribution is to embed this theoretical reason to herd in a model that is suitable to empirical analysis. When $\tau > 1$, extreme signals reveal the true value of the asset, and traders receiving them never herd. In the limit case of τ tending to infinity, all signal realizations become perfectly informative, with the result that no informed trader herds. Therefore, while our model allows for herd behavior, it also allows for the possibility that some traders (when $\tau > 1$) or all traders (when $\tau \rightarrow \infty$) rely only on their private information and never herd.

The probability of herding depends on the parameter values. To take an extreme example, when α (the probability of an information event) is arbitrarily close to zero, the market maker has a very strong prior that there is no information event. He barely updates the prices as trades arrive in the market, and herding arises as soon as there is an imbalance in the order flow, as happens in the seminal model of Bikhchandani, Hirshleifer, and Welch (1992). In contrast, if α is close to 1, the market maker and the informed traders update their beliefs in very similar manners, and herding rarely occurs.

Herding is important also for the informational efficiency of the market. During periods of herd behavior, private information is aggregated less efficiently by the

price as informed traders with good and bad signals may take the same action. The most extreme case is when traders herd for all signal realizations (e.g., traders herd buy even for $s_t^d = 0$). In such a case, the market maker is unable to make any inference on the signal realization from the trades. The market maker, however, updates his belief on the asset value, since the action remains informative on whether an information event has occurred.¹⁶ Since the market maker never stops learning, he gradually starts interpreting the history of past trades more and more similarly to the traders and, as a result, the measure of herders shrinks.

During an information-event day, the measure of herders changes with the sequence of trades and can become positive more than once at different times of the day. Given that information always flows to the market, however, the bid and ask prices converge to the asset value almost surely. Eventually the market maker learns whether a good event, a bad event, or no event occurred.¹⁷ Of course, with a finite number of trades, learning the true value of the asset is not guaranteed. As we mentioned above, an implicit assumption of the literature is that even in those days in which there is not enough trading activity, the market maker learns the true value of the asset before the following trading day starts (because public information is revealed during the night).¹⁸

II. The Likelihood Function

To estimate the herding model presented above, we have to specify its likelihood function. We write the likelihood function for the history of trades only, disregarding bid and ask prices.

In our model there is no public information: for this reason, there is a one-to-one mapping from trades to prices, and adding prices to the likelihood function would be redundant. The one-to-one mapping from trades to prices breaks down in the presence of public information, since price changes may be the result of public information arrival (as opposed to being only determined by the order flow). Nevertheless, our likelihood function for the history of trades would still be correctly specified, as long as the public news is independent of the informational event.

Remember also that information events are assumed to be independent and, as we mentioned in the previous section, before the market opens, market participants have learned the realization of the previous day's asset value. Because of this, the probability of the sequence of trades in a day depends only on the value of the asset that day. Therefore, the likelihood of a history of trades over multiple days can be written as

$$\mathcal{L}(\Phi; \{h^d\}_{d=1}^D) = \Pr(\{h^d\}_{d=1}^D | \Phi) = \prod_{d=1}^D \Pr(h^d | \Phi),$$

¹⁶The market maker learns since, in periods of herding, the proportion of buys and sells is different from that in an uninformed day. Essentially, whereas in our model there is herd behavior, there is never an informational cascade.

¹⁷The proof of convergence is standard, and we omit it. Recall that we have assumed that $(1 - \delta)\lambda^L = \delta\lambda^H$. This implies that $E(V_{d+1} | V_d = v_d) = v_d$. Since the price converges to the fundamental value almost surely, this guarantees that the martingale property of prices is satisfied. (We return to this point at the end of Section IV.)

¹⁸In any case, in our dataset there is enough trading activity that the true value of the asset is learned in most days: the end-of-day market maker's belief that an event has occurred is either above 0.9 or below 0.1 in 87 percent of days (i.e., in 87 percent of days the market maker has learned whether there was an event or not with 90 percent confidence).

where h^d is the history of trades at the end of a trading day d , and $\Phi := \{\alpha, \delta, \mu, \tau, \varepsilon\}$ is the vector of parameters.

Let us focus on the probability of a history of trades in a single day. As we have written, the *sequence* of trades, and not just the *number* of trades, conveys information. Having many buy orders at the beginning of the day is not equivalent to having the same number of buy orders spread out during the day. In fact, a particular sequence of buy or sell orders may create herd behavior: in periods of herding, the probability of a trade depends on the measure of informed traders who herd and is different from the probability in the absence of herding. Therefore, we have to compute the probability of a history of trades recursively, that is,

$$\Pr(h_t^d | \Phi) = \prod_{s=1}^t \Pr(x_s^d | h_s^d, \Phi),$$

where the probability of an action at time t of day d , $\Pr(x_t^d | h_t^d, \Phi)$, depends on the measure of informed traders who buy, sell, or do not trade after a given history of trades h_t^d . Using the law of total probability, at each time t , we compute $\Pr(x_t^d | h_t^d, \Phi)$ in the following way:

$$\begin{aligned} \Pr(x_t^d | h_t^d, \Phi) &= \Pr(x_t^d | h_t^d, v_d^H, \Phi) \Pr(v_d^H | h_t^d, \Phi) \\ &\quad + \Pr(x_t^d | h_t^d, v_d^L, \Phi) \Pr(v_d^L | h_t^d, \Phi) \\ &\quad + \Pr(x_t^d | h_t^d, v_{d-1}, \Phi) \Pr(v_{d-1} | h_t^d, \Phi). \end{aligned}$$

To show how to compute these probabilities, let us consider, first, the probability of an action conditional on a good-event day.¹⁹ For the sake of exposition, let us focus on the case in which the action is a buy order. As illustrated above, at each time t , in equilibrium there is a signal threshold β_t^d such that an informed trader buys for any signal realization greater than β_t^d , that is,

$$E(V_d | h_t^d, \beta_t^d) = a_t^d = E(V_d | h_t^d, X_t^d = \text{buy}, a_t^d, b_t^d),$$

which can be written as

$$\begin{aligned} v_{d-1} + \lambda^H \Pr(v_d^H | h_t^d, \beta_t^d) - \lambda^L \Pr(v_d^L | h_t^d, \beta_t^d) \\ = v_{d-1} + \lambda^H \Pr(v_d^H | h_t^d, \text{buy}_t^d) - \lambda^L \Pr(v_d^L | h_t^d, \text{buy}_t^d), \end{aligned}$$

or, after some manipulations, as²⁰

$$\begin{aligned} (1) \quad &\Pr(v_d^H | h_t^d, \beta_t^d) - \Pr(v_d^H | h_t^d, \text{buy}_t^d) \\ &= \frac{\delta}{1 - \delta} (\Pr(v_d^L | h_t^d, \beta_t^d) - \Pr(v_d^L | h_t^d, \text{buy}_t^d)). \end{aligned}$$

¹⁹For simplicity's sake, from now on, in the conditioning of the probabilities we will omit the vector of parameters Φ .

²⁰Note that, for simplicity's sake, in the probabilities to compute the ask we have omitted a_t^d and b_t^d in the conditioning. More importantly, note that the magnitude of the shocks that buffet the asset's value (λ^L and λ^H) do

The probabilities in this equation can easily be expressed as a function of the traders' and market maker's beliefs at time $t - 1$ and of the parameters. Specifically, the probability that an informed trader receiving signal β_t^d attaches to the good event is

$$\begin{aligned} & \Pr(v_d^H | h_t^d, \beta_t^d) \\ &= \frac{g^H(\beta_t^d | v_d^H) \Pr(v_d^H | h_t^d, V_d \neq v_{d-1})}{g^H(\beta_t^d | v_d^H) \Pr(v_d^H | h_t^d, V_d \neq v_{d-1}) + g^L(\beta_t^d | v_d^L) \Pr(v_d^L | h_t^d, V_d \neq v_{d-1})}. \end{aligned}$$

The probability that the market maker attaches to the good event can, instead, be computed as

$$\begin{aligned} & \Pr(v_d^H | h_t^d, buy_t^d) \\ &= \frac{\Pr(buy_t^d | v_d^H, h_t^d) \Pr(v_d^H | h_t^d)}{\Pr(buy_t^d | v_d^H, h_t^d) \Pr(v_d^H | h_t^d) + \Pr(buy_t^d | v_{d-1}^d, h_t^d) \Pr(v_{d-1}^d | h_t^d) + \Pr(buy_t^d | v_d^L, h_t^d) \Pr(v_d^L | h_t^d)}. \end{aligned}$$

By substituting these expressions into (1) we can compute β_t^d . Two comments are in order. First, the above expressions themselves contain the probabilities of a buy order by an informed trader in a good, bad, and no-event day; all these probabilities obviously depend on the threshold β_t^d itself (as illustrated below). That is, the threshold β_t^d is a fixed point. Second, at time $t = 1$, the prior beliefs of the traders and of the market maker are a function of the parameters only. Therefore, we can compute β_1^d as the solution to equation (1), and, from β_1^d , the probability of a buy order at time 1. After observing x_1^d , we update the market maker's and traders' beliefs, repeat the same procedure for time 2, and compute β_2^d and the probability $\Pr(buy_2^d | x_1^d, v_d^H)$. We do so recursively for each time t , always conditioning on the previous history of trades. Note that in order to maximize the likelihood function, the thresholds β_t^d (and the analogous threshold σ_t^d) have to be computed for each trading time in each day of trading, for each set of parameter values.

Once we have solved for β_t^d , we can compute the probability of a buy order in a good-event day. Let us focus on the case in which $\tau \in [0, 1]$; that is, let us concentrate on the case of bounded beliefs. In this case,

$$\begin{aligned} \Pr(buy_t^d | h_t^d, v_d^H) &= \mu \int_{\beta_t^d}^1 (1 + \tau(2s_t^d - 1)) ds_t^d + (1 - \mu) \left(\frac{\varepsilon}{2} \right) \\ &= \left(\tau(1 - \beta_t^{d2}) + (1 - \tau)(1 - \beta_t^d) \right) \mu + (1 - \mu) \left(\frac{\varepsilon}{2} \right). \end{aligned}$$

By following an analogous procedure, we compute σ_1^d and the probability of a sell order in a good-event day, that is,

$$\Pr(sell_t^d | h_t^d, v_d^H) = ((1 - \tau)\sigma_t^d + \tau\sigma_t^{d2})\mu + (1 - \mu) \left(\frac{\varepsilon}{2} \right).$$

The probability of a no-trade is just the complement to the probabilities of a buy and of a sell.

not appear in this equation, since the shocks cancel out. This is important, since it implies that we do not need to estimate them.

The analysis for the case of a bad information event ($V_d = v_d^L$) follows the same steps. The case of a no-event day ($V_d = v_{d-1}$) is easier, since the probability of a buy or of a sell is $\frac{\varepsilon}{2}$ and that of a no-trade is $1 - \varepsilon$. Also the case of unbounded beliefs, where $\tau \geq 1$, can be dealt with in a similar manner. The only changes are the extremes of integration when computing the probability of a trade.

Finally, to compute $\Pr(x_t^d | h_t^d, \Phi)$, we need the conditional probabilities of V_d given the history until time t , that is, $\Pr(V_d = v | h_t^d, \Phi)$ for $v = v_d^L, v_{d-1}, v_d^H$. These can also be computed recursively by using Bayes's rule.

To conclude, let us provide some intuition regarding the model's identification. For simplicity's sake, let us consider only the number of buys, sells, and no-trades in each day.²¹ Similarly to analogous structural models of market microstructure, our model classifies days into high-volume days with a prevalence of buys ("good-event" days), high-volume days with a prevalence of sells ("bad-event" days), and low-volume days ("no-event" days). The parameter α defines the probability that there is an event at the beginning of a trading day. We use data over many days of trading to identify it. The imbalance between buys and sells in event days identifies δ . No-event days allow us to identify ε , since in no-event days only noise traders trade. Finally, in good-event days, the ratio between buys and sells is determined by the proportion of traders who trade in the right direction (i.e., buy when there is a good event), which depends on μ and τ . An analogous argument holds for bad-event days. For any given estimate of μ and τ there is only one predicted ratio between buy and sell orders in the two types of days.²²

III. Data

Our purpose is to carry out a structural estimation of herding based on a market microstructure model. We perform our empirical analysis on a stock, Ashland Inc., traded on the New York Stock Exchange and already used in the seminal paper by Easley, Kiefer, and O'Hara (1997a).²³

We obtained the data from the TAQ (Trades and Quotes) dataset.²⁴ The dataset contains the posted bid and ask prices (the "quotes"), the prices at which the transactions occurred (the "trades"), and the time when the quotes were posted and when the transactions occurred. We used transactions data on Ashland Inc. in 1995, for a total of 252 trading days. The data refer to trading on the New York Stock Exchange, the American Stock Exchanges, and the consolidated regional exchanges.²⁵

The TAQ dataset does not sign the trades; that is, it does not report whether a transaction was a sale or a purchase. To classify a trade as a sell or a buy order, we used the standard algorithm proposed by Lee and Ready (1991). We compared the

²¹ In our estimation we use much more information than that, since we take into account the entire sequence of trades when constructing the likelihood function.

²² For a further argument for identification, see footnote 28.

²³ The name of the stock is slightly different, since the company changed name in 1995, and Easley, Kiefer, and O'Hara (1997a) use 1990 data.

²⁴ Hasbrouck (2004) provides a detailed description of this dataset.

²⁵ In the online Appendix, we present the results for seven more stocks, traded on the same exchanges. The estimates are quite similar to those reported for Ashland Inc., both in terms of parameter estimates (proportion of informed traders, frequency of information event, etc.) and in terms of occurrence of herding. We refer the reader to the online Appendix for further details.

transaction price with the quotes that were posted just before a trade occurred.²⁶ Every trade above the midpoint was classified as a buy order, and every trade below the midpoint was classified as a sell order; trades at the midpoint were classified as buy or sell orders according to whether the transaction price had increased (uptick) or decreased (downtick) with respect to the previous one. If there was no change in the transaction price, we looked at the previous price movement, and so on.²⁷

TAQ data do not contain any direct information on no-trades. We used the established convention of inserting no-trades between two transactions if the elapsed time between them exceeded a particular time interval (see, e.g., Easley, Kiefer, and O'Hara 1997a). We obtained this interval by computing the ratio between the total trading time in a day and the average number of buy and sell trades over the 252 days (see, e.g., Chung, Li, and McInish 2005). In our 252 trading-day window, the average number of trades per day was 90.2. We divided the total daily trading time (390 minutes) by 90.2 and obtained a unit-time interval of 259 seconds (i.e., on average, a trade occurred every 259 seconds). If there was no trading activity for 259 seconds or more, we inserted one or more no-trades to the sequence of buy and sell orders. The number of no-trades that we inserted between two consecutive transactions was equal to the number of 259-second time intervals between them. To check the robustness of our results, we also replicated the analysis for other no-trade time intervals (2, 3, 4, 5, 6, and 7 minutes).

Our sample of 252 trading days contained on average 149 decisions (buy, sell, or no-trade) per day. The sample was balanced, with 30 percent of buys, 31 percent of sells, and 40 percent of no-trades.

Finally, remember that in our theoretical model we assume that the closing price is a martingale. For the case of Ashland Inc. during 1995, the data support the hypothesis that the closing price is a martingale with respect to the history of past prices (i.e., the information available in our dataset): the autocorrelogram of price changes is not significantly different from zero, at all lags and at all significance levels (see the online Appendix).

IV. Results

We first present the estimates of the model parameters and then illustrate the importance of herd behavior in the trading activity of Ashland Inc. during 1995.

A. Estimates

We estimated the parameters through maximum likelihood, using both a direct search method (Nelder-Mead simplex) and the Genetic Algorithm.²⁸ The two

²⁶Given that transaction prices are reported with a delay, we followed Lee and Ready's (1991) suggestion of moving each quote ahead in time by five seconds. Moreover, following Hasbrouck (1991a), we ignore quotes posted by the regional exchanges.

²⁷We classified all trades with the exception of the opening trades, since these trades result from a trading mechanism (an auction) substantially different from the mechanism of trading during the day (which is the focus of our analysis).

²⁸We also simulated the theoretical model and verified that we could recover the model's parameters. Both methods converged to the true parameter values, which provides further evidence in favor of identification.

TABLE 1—ESTIMATION RESULTS

Parameter	Estimate	SD
α	0.28	0.03
δ	0.62	0.06
μ	0.42	0.01
τ	0.45	0.02
ε	0.57	0.002

Note: The table shows the estimates for the five parameters of the model and their standard deviations.

methods converged to the same parameter values. Table 1 presents the estimates and the standard deviations for the five parameters of the model.²⁹

Information events are relatively frequent: from the estimate of α , we infer that the probability of an information event is 28 percent, that is, in more than a fourth of trading days trading activity is motivated by private information. There is a small imbalance between good- and bad-event days: the probability of a good-information event is 62 percent (although the parameter δ has a relatively high standard deviation).³⁰ During event days, the proportion of traders with private information is 42 percent. The remaining trading activity comes from noise traders, who trade 57 percent of the time. Moreover, private information is noisy (that is, it is not perfectly informative). The estimate for τ is 0.45, which means that the probability of receiving an “incorrect signal”—i.e., a signal below 0.5 when we are in a good-event day or a signal above 0.5 when we are in a bad-event day—is 39 percent.³¹

As explained above, we constructed our dataset adding a no-trade after each 259 seconds of trading inactivity. As a robustness check, we repeated the estimation on several other datasets, where we added a no-trade for different intervals of trading inactivity. We report these estimates in Table 2.

The estimates of the probability of an information event (α) and of a good event (δ) are fairly similar over the different no-trade intervals. The estimate of ε increases with the no-trade interval: this is expected since the number of no-trades in the sample (and, therefore, also in the no-event days) becomes smaller and smaller. To have a characterization of trading activity in no-event days independent of the no-trade interval, following Easley, Kiefer, and O’Hara (1997a), we computed the probability of observing at least one trade during a five-minute interval in a no-event day: $\Lambda = 1 - (1 - \varepsilon)^{\frac{300}{\text{Seconds}}}$ (where “Seconds” is the no-trade interval). Table 2 shows this probability to be independent of the choice of the no-trade interval.

The parameter μ is quite stable across samples, whereas τ increases. To understand this, it is useful to observe that if both τ and μ were constant, as ε increases the estimated proportion of trading activity due to traders not having correct information

²⁹Standard deviations are computed numerically with the BHHH estimator.
³⁰Note that δ is greater than 0.5, although in the sample the number of buys and sells is essentially balanced. This happens because, among the days with high trading volume (classified as event days), a higher number of days have a positive trade imbalance than a negative one. To see this, consider the posterior beliefs of δ and α at the end of each day. In 22 percent of days the posterior belief of both α and δ is above 0.5 (i.e., a good-event day is more likely), whereas in only 12 percent of days the posterior belief of α is above 0.5 and that of δ is below 0.5 (i.e., a bad-event day is more likely).
³¹Given the signal density functions, the probability of an incorrect signal is given by $0.5 - 0.25\tau$.

TABLE 2—ROBUSTNESS CHECKS FOR DIFFERENT NO-TRADE INTERVALS

	NT = 120	SD	NT = 180	SD	NT = 240	SD
α	0.20	0.02	0.25	0.02	0.27	0.02
δ	0.72	0.02	0.67	0.04	0.61	0.01
μ	0.26	0.01	0.36	0.01	0.41	0.01
τ	0.44	0.01	0.40	0.03	0.44	0.02
ε	0.33	0.001	0.45	0.001	0.54	0.002
Λ	0.64		0.63		0.62	
Γ	0.32		0.33		0.34	
	NT = 300	SD	NT = 360	SD	NT = 420	SD
α	0.30	0.03	0.27	0.03	0.27	0.04
δ	0.60	0.03	0.61	0.03	0.55	0.07
μ	0.40	0.01	0.39	0.01	0.37	0.01
τ	0.52	0.03	0.59	0.03	0.67	0.01
ε	0.62	0.002	0.69	0.002	0.74	0.002
Λ	0.62		0.62		0.62	
Γ	0.32		0.31		0.30	

Notes: The table shows the estimates for various no-trade intervals, from two to seven minutes. The last two rows report two more statistics derived from the estimated parameters and explained in the text.

(either because they are noise or because their signal is incorrect) would increase. In contrast, this proportion should obviously be independent of our choice of no-trade interval. This is indeed the case. To show this we computed the parameter

$$\Gamma = \frac{\mu(0.5 + 0.25\tau)}{(1 - \mu)\varepsilon + \mu},$$

which represents the proportion of correctly informed traders (e.g., informed traders with a signal greater than 0.5 in a good-event day) over the sum of all informed traders and the noise traders who trade. In other words, Γ is approximately equal to the fraction of trades coming from informed traders with the correct signal.³² It is remarkable that Γ , which equals 0.34 when the no-trade interval is 259 seconds, is constant across all the different datasets that we used to estimate the model’s parameters. This shows the robustness of our results to the choice of the no-trade interval.

Let us now discuss how our results compare to different specifications of the model. A natural comparison is with a model in which the signal precision is not estimated but is restricted to be perfectly informative (i.e., $\tau \rightarrow \infty$). This is the case studied by Easley, Kiefer, and O’Hara (1997a). In this case, all informed traders follow their own private information, the sequence of trades has no informational content beyond the aggregate numbers of buys, sells, and no-trades, and herding never arises. As a result, the likelihood function does not need to be computed recursively (see Easley, Kiefer, and O’Hara 1997a, for a detailed description). Table 3 presents the estimated parameters.

The estimates for α and δ are very close to those we obtained for our model. This shows that the classification of days is not affected by the specification of the signal

³²The approximation is due to the fact that because of the bid-ask spread, we are ignoring that a small measure of informed traders may not trade. Easley, Kiefer, and O’Hara (1997a) report a similar composite parameter when analyzing their results for different no-trade intervals.

TABLE 3—PARAMETER ESTIMATES FOR THE EASLEY, KIEFER, AND O’HARA (1997a) MODEL

Parameter	Estimate	SD
α	0.33	0.04
δ	0.60	0.06
μ	0.17	0.01
ε	0.58	0.002

Notes: The table shows the estimates for the four-parameter model of Easley, Kiefer, and O’Hara (1997a), in which informed traders know the true asset value. The no-trade interval is 259 seconds.

structure. Similarly, the estimates for ε in the two models are almost the same. This is not surprising, since ε captures the trading activity of noise traders and is not affected by assumptions on the structure of private information. The parameter μ is smaller in the restricted model, which is intuitive since this model imposes that all informed traders receive the correct signal (i.e., they know whether a good- or a bad-information event occurred).

The restriction in Easley, Kiefer, and O’Hara (1997a) is not supported by the data. The likelihood ratio test overwhelmingly rejects the restriction of perfectly informative signals, with a LR statistic of 272.15 (and a p -value of zero).³³ This is important for our aims, since the fact that signals are not perfectly informative implies that the sequence in the order flow matters. In other words, the number of buys, sells, and no-trades at the end of the day is not a sufficient statistic for the pattern of trading activity. Depending on the sequence, herd behavior by informed traders may occur in equilibrium.³⁴

In the market microstructure literature, a great deal of attention has been given to the PIN, a measure of the probability that a trade comes from an informed trader (see, among others, Easley et al. 1996, and the literature cited in Chung, Li, and McNish 2005). This measure is given by
$$PIN = \frac{\alpha\mu}{\alpha\mu + \varepsilon(1 - \alpha\mu)},$$
 where the numerator is the beginning-of-the-day probability that a trade is information based and the denominator is the probability that a trade occurs. With the estimated parameters of our model, the PIN equals 19 percent, whereas for the Easley, Kiefer, and O’Hara (1997a) model it is only 9 percent.³⁵ The difference is due to the fact that, in the previous literature, incorrect trades (e.g., sell orders in a good-event day) can only be due to exogenous, noninformative (e.g., liquidity) reasons to trade, whereas in our setup we do not exclude that they may come from informed traders who either receive the incorrect information or herd. Because of this, the PIN computed for the Easley, Kiefer, and O’Hara (1997a) model is lower than for our model.

³³That is, we reject the null that the signal is perfectly informative; the log-likelihood ratio of the restricted and unrestricted models, once appropriately scaled, is distributed as a χ^2 with one degree of freedom. A note of warning on the result of the test is needed here, since the null hypothesis is on the boundary of the parameter space (see Andrews 2001).

³⁴In the online Appendix, we describe the pattern of the data that produces the statistical rejection of a model with perfect information. In particular, in such a model, the probability of buy and sell orders in event days should be constant. We show with a logit regression analysis that this is not the case.

³⁵We compute the PIN for our model using the same formula as Easley et al. (1996). They interpret the PIN as the probability of a trade coming from an informed trader at the beginning of the day. In our model, since the signal is continuous, the interpretation is correct only if we ignore the bid-ask spread (otherwise, some informed traders may decide not to trade because their expectations fall inside the bid-ask spread). We use this approximation for simplicity’s sake and to keep comparability with the existing work on the PIN.

If we adjust for the fact that in our model the information may not be correct (i.e., we multiply the PIN computed with our parameter estimates by the probability of a correct signal $0.5 + 0.25\tau$), the proportion of trading activity stemming from traders with a *correct signal* becomes almost the same as the standard PIN in Easley, Kiefer and O'Hara's (1997a) model. Since the null that the signal is perfectly precise is rejected by the data, our results suggest that the PIN computed from a model with signals that are always correct (that is, as computed in the literature) measures the proportion of informed-based trading coming from traders *receiving the correct information* and not the overall proportion of information-based trading (which is its usual interpretation).

Finally, note that a 99 percent confidence interval for τ does not include 1.³⁶ This means that there is evidence in our sample that there are no realizations of the signal that reveal the true asset value with probability one. In the jargon of the social learning literature, signals are bounded.

B. Herd Behavior

The estimates of the parameters α and τ imply that herd behavior can occur in our sample. Since the estimate of α is clearly lower than 1,³⁷ there is event uncertainty in the market, which is a necessary condition for the mechanism of herd behavior highlighted in Section IA to work. Moreover, the estimate $\tau = 0.44$ means that traders receive a signal that is noisy (i.e., not perfectly informative) and may decide to act against it (i.e., buy upon receiving a bad signal or sell upon receiving a good one).

The Frequency of Herding.—Recall that there is herd behavior at time t of day d when there is a positive measure of signal realizations for which an informed trader either herd buys or herd sells, that is, when, in equilibrium, either $\beta_t^d < 0.5$ (herd buy) or $\sigma_t^d > 0.5$ (herd sell). To gauge the frequency of herd behavior in our sample, for each trading day we computed the buy thresholds (β_t^d) and the sell thresholds (σ_t^d) given our parameter estimates. As an illustration, Figure 5 shows the thresholds (on the right vertical axis) for one day out of the 252 days in the sample. Whenever the buy threshold (dashed line) drops below 0.5 or the sell threshold (solid line) goes above 0.5, there is herd behavior. The shaded area (measured on the left vertical axis) represents the trade imbalance, that is, at each time t , the number of buys minus the number of sells arrived in the market from the beginning of the day until time $t - 1$. As one can see, herd buying occurs at the beginning of trading activity, as the trade imbalance is positive, that is, as more buys than sells arrive in the market. This is followed by a long stretch of herd sells, as sell orders arrive, and the trade imbalance becomes negative. At the very end of the day, herd behavior effectively disappears.

To understand better the informed traders' behavior, it is useful to look at how the market maker changes his expectation and sets the prices during the day. Figure 6 reports the evolution of the price (i.e., the market maker's expectation) during the day and the probability that the market maker attaches to being in an event

³⁶Since the parameter's standard deviation is 0.02, this is the case for any reasonable confidence interval.

³⁷The parameter's standard deviation is 0.03. See the argument in the previous footnote.




FIGURE 5. A DAY OF TRADING

Notes: The figure reports the evolution of the trade imbalance (shaded line), buy threshold (dashed line), and sell threshold (solid line) in one day of trading. The thresholds are measured on the right vertical axis and the trade imbalance on the left vertical axis. Herd selling occurs when the solid line is above 0.5 (indicated by a horizontal line) and herd buying when the dashed line falls below 0.5.

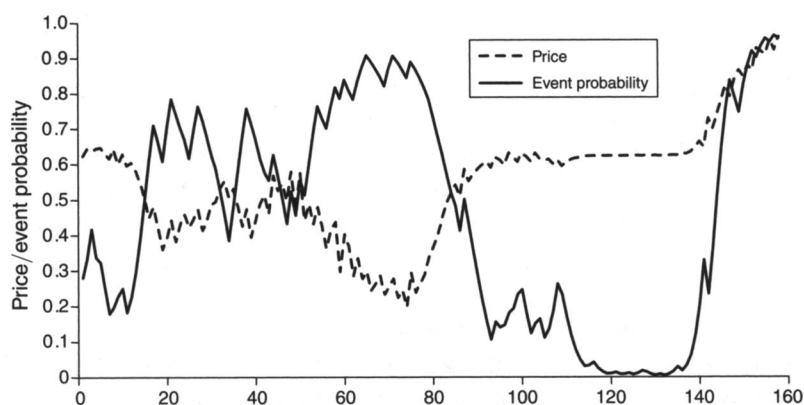


FIGURE 6. A DAY OF TRADING

Note: The figure reports the evolution of the price (dashed line) and of the probability that the market maker attaches to being in an event day (solid line) in one day of trading.

day. For the first 100 periods, the market maker's belief on the occurrence of an event fluctuates because, although the sell orders outnumber the buy orders, there are nevertheless many buy orders and several periods of inactivity. After trading time 90, the market maker is confident that there was no event, as his belief on the occurrence of an information event approaches zero.³⁸ As a result, he does not need to protect himself from adverse selection, and, as illustrated above, the buy and sell thresholds become very close to each other. Finally, the sharp increase in the trade imbalance after period 100 leads the market maker to reassess (i.e., eventually, to

³⁸The sharp reduction in the probability the market maker attaches to an information event between periods 80 and 100 is due to the fact that although we are in a period of herd sell, the trade imbalance is fairly constant.

TABLE 4—THE FREQUENCY OF HERDING

	> 10 percent	> 30 percent	> 50 percent	Mean time	Mean length	Length SD	Max length
Herd buy	0.79	0.47	0.23	60	9	13	99
Herd sell	0.83	0.58	0.35	61	10	15	120

Notes: The left part of the table shows the proportion of days in which the percentage of trading periods with herd behavior was higher than 10, 30, or 50. For instance, in 23 percent of days, herd-buy periods were at least 50 percent of the total. The right part of the table reports some herding statistics. The first column shows the mean trading period in which herd behavior occurred. The other columns show the mean, standard deviation, and maximum for the number of consecutive trading periods in which there was herd behavior.

increase) the probability he attaches to an information event, thus leading to a widening of the thresholds.³⁹

In order to analyze the frequency of herding across all the days in our sample, we compute how often the buy threshold is below 0.5 or the sell threshold is above 0.5. It should be clear that this analysis is relevant both for event and no-event days. In both cases, the existence of signal realizations for which informed traders herd (if an information event has occurred, which neither the market maker nor an external observer knows at the moment of the trade) modifies the way the market maker updates the price, which, as we shall see, affects the market’s informational efficiency.

In our sample, herding happens quite frequently: over the 252 days of trading, β_t^d is below 0.5 in 30 percent of trading periods, and σ_t^d is above 0.5 in 37 percent of trading periods. Moreover, there are some days where herding is very pronounced. The left part of Table 4 reports the proportion of days in which the buy or sell thresholds (i.e., β_t^d or σ_t^d) cross 0.5 at least 10, 30, and 50 percent of the time. Herd buying was observed in more than 50 percent of the trading times in 23 percent of the 252 days in our sample. Similarly, herd selling was observed in more than 50 percent of the trading times in 35 percent of the 252 days in our sample.

The right part of Table 4 reports the mean period of the day when we observed herding, that is, approximately the sixtieth trading period for both herd buying and herd selling (which, in clock time, is roughly after 2 hours and 36 minutes of trading).⁴⁰ It is also interesting to ask how long periods of herding last, that is, for how many trading periods we observe a positive measure of herders after a herd starts. Herd buys last on average 9 trading periods (corresponding to about 24 minutes), and herd sells last on average for 10 periods. There is, however, pronounced variability in the length of herds, with a standard deviation of 13 and 15 trading periods. The longest herd buy lasted for 99 periods (about 257 minutes) and the longest herd sell lasted 120 periods (312 minutes).

Proportion of Herders.—The previous analysis helps us gauge how often herding occurs in our sample. The fact that at a given time the buy (sell) threshold is lower (higher) than 0.5, however, does not tell us how likely it is for an informed trader to herd at that time. This is captured by the measure of signal realizations for

³⁹As one can observe in Figure 6, the price evolution is quite smooth, whereas the market maker’s belief on the occurrence of an event swings considerably. This is due to the extreme asymmetry of information between traders and market makers as to whether an event has occurred.
⁴⁰On average there are 149 decisions in a day (390 minutes). Therefore, the average length of a trading period is 2.6 minutes.

TABLE 5—PERCENTAGE OF HERDERS

	Average	SD	Max	> 1 percent	> 5 percent	> 10 percent
Herd buyers	2 percent	3 percent	11 percent	0.53	0.16	0.07
Herd sellers	4 percent	6 percent	29 percent	0.55	0.22	0.11

Notes: The left part of the table shows the mean, the standard deviation, and the maximum of the percentage of herd buyers and herd sellers; the right part of the table shows the proportion of days in which the percentage of herd buyers or herd sellers was higher than 1, 5, or 10. For instance, in 11 percent of days, the percentage of herd sellers was at least 10.

which an informed trader herds. As we will discuss in detail in the next subsection, this measure is also very important for the informational efficiency of the market. The higher the measure of signal realizations for which traders herd, the lower the informational efficiency.

To compute the measure of herd-signal realizations, however, we need to know the distribution of signals on any given day. This, in turn, depends on whether the day of trading was a good-event or a bad-event day. To this purpose, we classified a day as a good-event day (bad-event day) if either $\Pr(V_d = v^H | h_{T_d}) > 0.9$ (good-event day) or $\Pr(V_d = v^L | h_{T_d}) > 0.9$ (bad-event day). That is, we classified a day as a good- (bad-) event day if, at the end of the day, the posterior probability of the day being a good- (bad-) event day was higher than 0.9.⁴¹

We concentrated our analysis on the days classified as good-event or bad-event days. For each trading period, we computed the proportion of bad signals for which informed traders would herd buy and the proportion of good signals for which informed traders would herd sell. Consider, for instance, a good-event day. On such a day, the signal is distributed according to $g^H(s_t^d | v_d^H) = 1 + 0.44(2s_t^d - 1)$. For each trading period in which $\beta_t^d < 0.5$, we computed the measure of signals between β_t^d and 0.5 (i.e., the measure of signals for which an informed trader herd buys). We then divided this measure by the measure of signals between 0 and 0.5 (the measure of all bad signals, i.e., all signals for which the informed trader could potentially herd buy). We refer to this ratio as the “proportion of herd buyers,” since it is the proportion of informed traders with a bad signal who would nevertheless herd buy were they to trade in that period. The proportion of herd sellers was calculated in a similar way. We report the average results in Table 5.

As the left part of Table 5 shows, on average, across all event days, the proportion of herd buyers was 2 percent and that of herd sellers 4 percent. These proportions are, however, quite variable across days, reaching a maximum of 11 percent for herd buying and 29 percent for herd selling. Misdirected herding (i.e., herd buying in a bad-event day and herd selling in a good-event day) does occur: on average, in a bad-event day, the proportion of herd buyers was 1 percent; in a good-event day, the proportion of herd sellers was 2 percent.

As the last three columns of Table 5 show, there are a substantial number of days where the percentage of herd buyers or sellers is significant: for instance, in 7 percent of event days, the proportion of informed traders who herd buy was higher than

⁴¹Of course, the 0.9 threshold is arbitrary. As a robustness check, we repeated the calculations for 0.75, 0.8, and 0.85 and obtained very similar results. The results are available in the online Appendix.

TABLE 6—PROPORTION OF BUYS AND SELLS IN PERIODS OF HERDING

	Good-event days	Bad-event days
$\Pr(\text{buy} \beta_t^d < 0.5)$	0.48	0.31
$\Pr(\text{buy})$	0.43	0.26
$\Pr(\text{sell} \sigma_t^d > 0.5)$	0.39	0.48
$\Pr(\text{sell})$	0.27	0.45

Note: The table shows the proportion of buys and sells in periods of herd buy, in periods of herd sell, and in the whole sample.

10 percent; similarly, in 11 percent of event days, the proportion of informed traders who herd sell was higher than 10 percent. This confirms the result of the previous section that herding behavior seems to be particularly concentrated in some days of trading.

An important question is whether herding usually happens after buy or sell orders have accumulated in the market. We computed the average level of the trade imbalance in periods of herd buy (i.e., when $\beta_t^d < 0.5$). The trade imbalance is on average positive both in good- and in bad-event days (9 and 3.5, respectively). This means that herd buying usually happens when there has been a preponderance of buys. Similarly, herd selling usually occurs when there has been a preponderance of sells (the trade imbalance is -8.2 in good-event days and -12.9 in bad-event days).

By definition, herd buying increases the proportion of buys, and herd selling increases the proportion of sells. Table 6 illustrates this point by showing the frequencies of buy orders and of sell orders that we observe in periods of herd buying and herd selling and contrasting them with the overall frequencies. In good-event days, for instance, the overall frequency of buy orders is 43 percent. This frequency goes up to 48 percent when there is herd buying. Two comments are in order. First, the pattern shown in Table 6 is inconsistent with models in which the signals are always correct (as in Easley, Kiefer, and O’Hara 1997a). In this class of models, the probability of any action is constant as long as one conditions on the type of day (good-event day, bad-event day, or no-information day). This is not true in the data and is a deviation from the prediction of Easley, Kiefer, and O’Hara (1997a) that our analysis captures (since, in our model, the probability of a given action changes with the whole sequence of trading activity, as the buy and the sell thresholds change). Second, since as mentioned above, herding usually happens after buy or sell orders have accumulated in the market, these results imply that higher positive (negative) levels of the trade imbalance increase the probability of a buy (sell) order. That is, herd behavior generates serial dependence in the trading pattern during the day.⁴²

⁴² In a related strand of literature, Hasbrouck (1991a, 1991b) and Chung, Li, and McNish (2005), among others, provide evidence of autocorrelation in trades. Easley, Kiefer, and O’Hara (1997b) and Easley et al. (2008) recognize the importance of serial dependence in trading activity. To capture it, they allow for path dependence in noise trading due to unmodeled reasons. In contrast, in our model, the sequence of trading becomes important for reasons dictated by economic theory.

Nevertheless, one may wonder whether the presence of herding hinges on the assumption that the arrival of noise traders is an i.i.d. process; i.e., whether what we call herding is instead a time dependence in liquidity shocks that we assume away in our maximum likelihood estimation. This is not the case. In the online Appendix, we report the results of the estimation of alternative model specifications in which we allow for time-dependence in the behavior of noise traders. We find that our estimate of the amount of herding in the market does not decrease; if anything, depending on the specific form of time dependence, herding may slightly increase.

As a final remark, let us observe that our analysis of herding is very different from the existing, statistical (i.e., atheoretical) measures of herding. As we mentioned in the Introduction, these measures of herding (starting with the paper of Lakonishok, Shleifer, and Vishny 1992—henceforth, LSV) do not study informational herd behavior in the trading activity of a security. Instead, they focus on whether portfolio or trading decisions by investors are clustered. This action clustering, labeled “herding” by this literature, may derive from real imitation, but also from many other factors (e.g., common reaction to a public announcement). Additionally, the LSV measure has been traditionally used on portfolio holdings of fund managers at the monthly or quarterly frequency, rather than on transaction data. For these reasons, a comparison between the two types of analyses is not straightforward. Nevertheless, to offer a point of reference, in the online Appendix we followed a recent paper by Christoffersen and Tang (2010) and computed the LSV measure of herding using intraday transaction data. For Ashland Inc. in 1995 the LSV measure of herding is approximately 4 percent: that is, on average, in a day, there are 4 percent buy (or sell) orders on Ashland Inc. in excess of what we would observe in the absence of action clustering (i.e., of herding as defined by LSV). In order to compare the LSV measure to our results, we computed an “ML-based” LSV measure, which we defined as the level of LSV that is implied by our maximum-likelihood estimates of informational herding. This “ML-based” LSV measure is only 1 percent (i.e., on average, in a day, there are 1 percent buy (or sell) orders on Ashland Inc. in excess of what we would observe in the absence of informational herding). We conclude that the traditional LSV methodology, when applied to transaction data, classifies a much higher proportion of transactions as being due to herding than the proportion actually due to informational herding (as implied by our structural estimates). The reason is simple. Consider, for instance, event days. Informed traders cluster their decisions because they react to the same piece of private information (e.g., in a good-event day, 61 percent of informed traders receive a good signal); the LSV methodology cannot disentangle their reaction to common information from true informational herding.⁴³ We refer the reader to the online Appendix for a detailed analysis.

Informational Inefficiency.—In periods of herd behavior, a proportion of informed traders trade against their signals; as a result, information is aggregated less efficiently by the asset price. Indeed, it is easy to show that trades convey the maximal amount of information when informed traders buy upon receiving a good signal and sell upon receiving a bad one.⁴⁴ In periods of herding, in contrast, traders may buy even with a bad signal or sell even with a good one.

To quantify the informational inefficiency caused by herding, we proceeded in the following way. We simulated the history of trades and prices over many days for our theoretical model, using our estimates of the parameter values. We then compared the simulation results with two benchmarks that capture the price behavior in an

⁴³This limitation of the LSV measure may be much less of an issue for the datasets to which it is usually applied (i.e., portfolio holdings by fund managers at the monthly or quarterly frequency).

⁴⁴By this, we mean that the strategy $X_t^d = -1$ for $s_t^d < 0.5$ and $X_t^d = 1$ for $s_t^d > 0.5$ minimizes $E[(E(V_d | h_t^d, x_t^d) - V_d)^2 | h_t^d]$. Note that, for simplicity's sake, we assume that traders cannot use no-trades to reveal information, since in equilibrium they abstain from trading only when their expectation is inside the bid-ask spread.

TABLE 7—INFORMATIONAL INEFFICIENCY

No-herd benchmark		No-event uncertainty benchmark
All days	Event days	Event days
4 percent	10 percent	12 percent

Note: The table shows the average distance between the public belief in the model and that in the benchmarks as a percentage of the asset’s expected value.

informationally efficient market. In the first benchmark, we simulated the model forcing informed traders to buy (sell) upon a good (bad) signal. In other words, informed traders (irrationally) never herded and always followed their private information. As a second benchmark, we considered the case in which there is no event uncertainty, that is, the market maker knows whether there has been an informational event. As a result, he updates his beliefs (and prices) exactly as informed traders do, and, because of this, informed traders never herd. Essentially, in the first benchmark, the market is efficient because traders (irrationally) follow their signals; in the second benchmark, the market is efficient because the informational asymmetry between traders and market makers due to event uncertainty is eliminated.

We simulated the price paths for 1,000,000 days of trading (with 149 trading periods per day) for our theoretical model and for the two benchmarks. Then, at each time t of any day d , we computed the distance (i.e., the absolute value of the difference) between the price $E(V_d|h_t^d)$ in our model and that in each benchmark.

Table 7 presents the average distance taken over all trading periods as a percentage of the expected value of the asset. For the first benchmark, we present the average distance both over all days and over event days only. For the second benchmark, by construction, only the average distance over event days is meaningful.

The average distance between the prices in our model and in the first benchmark over all days amounts to 4 percent of the asset’s expected value. If we focus our attention on event days, the distance is higher (10 percent), since the imbalance between buys and sells causes herding to occur more often. The average distance between the public beliefs using the second benchmark is similar, equal to 12 percent. We also repeated these computations for days in which herding was more pronounced. Considering, for instance, the first benchmark, we find that in 7 percent of days (20 percent of event days) the distance between the price and the nonherding price is greater than 15 percent; similarly, in 4 percent of days (11 percent of event days) the distance between the price and the nonherding price is greater than 20 percent. This suggests that there are days when intraday herding affects the informational properties of the price in a very significant manner.^{45,46}

⁴⁵In order to give the reader a point of reference, let us note that the daily standard deviation of Ashland Inc.’s price in 1995 was 1 percent.

⁴⁶Note that by construction any mispricing due to herding is corrected by the beginning of the following day (by which point the previous day’s asset value is learned). This, however, does not lower the relevance of such mispricings. Trades are executed throughout the day at prices that are not the closing prices. Inasmuch as there are deviations from the informationally efficient price, this will affect the welfare of the traders, including those who are not in the market for speculative reasons.

V. Conclusion

We developed a theoretical model of herd behavior in financial markets that is amenable to structural estimation with transaction data. We estimated the model using data for a NYSE stock (Ashland Inc.) in 1995 and identified the periods in each trading day when informed traders herd. We found that herding was present in the market and fairly pervasive on some trading days. Moreover, herding generated important informational inefficiencies.

The main contribution of this article is methodological: it provides an empirical strategy to analyze herding within a structural estimation framework. This contrasts with existing empirical studies of herding, which are based on atheoretical, statistical measures of trade clustering.

In future research, we plan to use our methodology to investigate the importance of herding for a large number of stocks, by analyzing how herding changes with stock characteristics (e.g., large stocks versus small ones) and with the macroeconomic environment (e.g., crises versus tranquil periods). We also plan to contribute to the existing literature on information and asset pricing, both by studying whether our measure of market informativeness (which takes into account that information may not be perfectly precise) improves the performance of the information factor and by testing whether herding itself is a risk factor priced in the market.

Finally, whereas our interest was in learning in financial markets, our methodology could be fruitfully used in fields other than financial economics. The voluminous and growing theoretical literature on social learning has been intensively tested in laboratory experiments. It is, however, generally acknowledged that these models cannot be easily studied with field data because we lack data on private information. Our methodology shows how this problem can be overcome.

REFERENCES

- Andrews, Donald W. K. 2001. "Testing When a Parameter Is on the Boundary of the Maintained Hypothesis." *Econometrica* 69 (3): 683–734.
- Avery, Christopher, and Peter Zemsky. 1998. "Multidimensional Uncertainty and Herd Behavior in Financial Markets." *American Economic Review* 88 (4): 724–48.
- Banerjee, Abhijit V. 1992. "A Simple Model of Herd Behavior." *Quarterly Journal of Economics* 107 (3): 797–817.
- Bikhchandani, Sushil, David Hirshleifer, and Ivo Welch. 1992. "A Theory of Fads, Fashion, Custom, and Cultural Change in Informational Cascades." *Journal of Political Economy* 100 (5): 992–1026.
- Bikhchandani, Sushil, and Sunil Sharma. 2000. "Herd Behavior in Financial Markets." *IMF Staff Papers* 47 (3): 279–310.
- Christoffersen, Susan, and Ya Tang. 2010. "Institutional Herding and Information Cascades: Evidence from Daily Trades." Social Science Research Network, Working Paper 1572726.
- Chung, Kee H., Mingsheng Li, and Thomas H. McNish. 2005. "Information-Based Trading, Price Impact of Trades, and Trade Autocorrelation." *Journal of Banking and Finance* 29 (7): 1645–69.
- Cipriani, Marco, and Antonio Guarino. 2005. "Herd Behavior in a Laboratory Financial Market." *American Economic Review* 95 (5): 1427–43.
- Cipriani, Marco, and Antonio Guarino. 2008. "Herd Behavior and Contagion in Financial Markets." *B E Journal of Theoretical Economics: Contributions to Theoretical Economics* 8: Article 24.
- Cipriani, Marco, and Antonio Guarino. 2009. "Herd Behavior in Financial Markets: An Experiment with Financial Market Professionals." *Journal of the European Economic Association* 7 (1): 206–33.
- Cipriani, Marco, and Antonio Guarino. 2014. "Estimating a Structural Model of Herd Behavior in Financial Markets: Dataset." *American Economic Review*. <http://dx.doi.org/10.1257/aer.104.1.224>.

- Dasgupta, Amil, Andrea Prat, and Michela Verardo. 2011. "Institutional Trade Persistence and Long-Term Equity Returns." *Journal of Finance* 66 (2): 635–53.
- Drehmann, Mathias, Jörg Oechssler, and Andreas Roider. 2005. "Herding and Contrarian Behavior in Financial Markets: An Internet Experiment." *American Economic Review* 95 (5): 1403–26.
- Easley, David, Robert F. Engle, Maureen O'Hara, and Liuren Wu. 2008. "Time-Varying Arrival Rates of Informed and Uninformed Trades." *Journal of Financial Econometrics* 6 (2): 171–207.
- Easley, David, Nicholas M. Kiefer, and Maureen O'Hara. 1997a. "One Day in the Life of a Very Common Stock." *Review of Financial Studies* 10 (3): 805–35.
- Easley, David, Nicholas M. Kiefer, and Maureen O'Hara. 1997b. "The Information Content of the Trading Process." *Journal of Empirical Finance* 4 (2–3): 159–86.
- Easley, David, Nicholas M. Kiefer, and Maureen O'Hara, and Joe Paperman. 1996. "Liquidity, Information, and Infrequently Traded Stocks." *Journal of Finance* 51 (4): 1405–36.
- Easley, David, and Maureen O'Hara. 1987. "Price, Trade Size, and Information in Securities Markets." *Journal of Financial Economics* 19 (1): 69–90.
- Easley, David, and Maureen O'Hara. 1992. "Time and the Process of Security Price Adjustment." *Journal of Finance* 47 (2): 576–605.
- Gale, Douglas. 1996. "What Have We Learned from Social Learning?" *European Economic Review* 40 (3–5): 617–28.
- Gervais, Simon. 1997. "Market Microstructure with Uncertain Information Precision: A Multi-Period Analysis." Unpublished.
- Glosten, Lawrence R., and Paul R. Milgrom. 1985. "Bid, Ask, and Transaction Prices in a Specialist Market with Heterogeneously Informed Traders." *Journal of Financial Economics* 14 (1): 71–100.
- Grinblatt, Mark, Sheridan Titman, and Russ Wermers. 1995. "Momentum Investment Strategies, Portfolio Performance, and Herding: A Study of Mutual Fund Behavior." *American Economic Review* 85 (5): 1088–105.
- Hasbrouck, Joel. 1991a. "The Summary Informativeness of Stock Trades: An Econometric Analysis." *Review of Financial Studies* 4 (3): 571–95.
- Hasbrouck, Joel. 1991b. "Measuring the Information Content of Stock Trades." *Journal of Finance* 46 (1): 179–207.
- Hasbrouck, Joel. 2004. "Working with the WRDS (TAQ and CRSP) Data." Unpublished.
- Hirshleifer, David, and Siew Teoh. 2009. "Thought and Behavior Contagion in Capital Markets." In *Handbook of Financial Markets*, edited by Klaus Reiner Schenk-Hoppé and Thorsten Hens, 1–56. Amsterdam: Elsevier, North-Holland.
- Lakonishok, Josef, Andrei Shleifer, and Robert W. Vishny. 1992. "The Impact of Institutional Trading on Stock Prices." *Journal of Financial Economics* 32 (1): 23–43.
- Lee, Charles M. C., and Mark J. Ready. 1991. "Inferring Trade Direction from Intraday Data." *Journal of Finance* 46 (2): 733–46.
- Lee, In Ho. 1998. "Market Crashes and Informational Avalanches." *Review of Economic Studies* 65 (4): 741–59.
- NYSE. 1995. "TAQ dataset." Wharton Research Data Services (WRDS), <http://whartonwrds.com/archive-pages/our-datasets/nyse-euronext-taq/> (accessed January 2005).
- Park, Andreas, and Hamid Sabourian. 2011. "Herding and Contrarian Behavior in Financial Markets." *Econometrica* 79 (4): 973–1026.
- Smith, Lones, and Peter Sørensen. 2000. "Pathological Outcomes of Observational Learning." *Econometrica* 68 (2): 371–98.
- Vives, Xavier. 2008. *Information and Learning in Markets: The Impact of Market Microstructure*. Princeton: Princeton University Press.
- Welch, Ivo. 1992. "Sequential Sales, Learning, and Cascades." *Journal of Finance* 47 (2): 695–732.
- Wermers, Russ. 1999. "Mutual Fund Herding and the Impact on Stock Prices." *Journal of Finance* 54 (2): 581–622.